

Reward-Conditioned Reverse Diffusion for Portfolio Optimization

Siddhant Sukhani

Department of Management Science & Engineering

Stanford University

sukhani@stanford.edu

*MS&E 342: Stochastic Control & Optimization in Continuous Time
Spring Quarter 2026*

May 24, 2026

Abstract

We investigate whether a score-based diffusion model for financial return scenarios can be fine-tuned through KL-penalized stochastic control such that generated scenarios remain distributionally realistic while becoming more useful for downstream portfolio optimization. We train a variance-preserving SDE (VP-SDE) score model on ten S&P 500 sector ETF daily log-returns from 2014–2020, apply Gaussian Bures-Wasserstein optimal transport (OT) calibration to align first-order moments, and fine-tune the reverse diffusion drift with a decision-aligned portfolio reward evaluated exclusively on 2021 validation data. The fine-tuning objective balances expected portfolio utility against a Girsanov KL divergence penalty, with the regularization strength η (a positive multiplier on the KL penalty) selected on the validation period to prevent test-period leakage. Over the 2022–2024 test window, the fine-tuned strategy achieves a Sharpe ratio of **0.70** versus 0.52 for rolling plug-in Markowitz and 0.36 for equal-weight, with a reduced maximum drawdown of 14.9%. A formal leakage audit confirms zero data-split violations. Bootstrap confidence intervals (95%) span approximately 2.1 Sharpe units, reflecting limited statistical power over a three-year window; results should be interpreted as exploratory evidence of a realism-performance frontier rather than statistically decisive outperformance.

1 Introduction

Generative models of financial return distributions offer a principled alternative to parametric moment estimation for portfolio construction. Rather than fitting a multivariate Gaussian and applying Markowitz mean-variance optimization (MVO) (Markowitz, 1952), one can sample *scenarios* from a learned generative model and solve a sample-average stochastic program. This scenario-based approach can capture non-Gaussian features that are empirically prominent in financial data, including fat tails, left skewness, and volatility clustering (Cont, 2001; McNeil et al., 2015).

Score-based diffusion models have become the dominant class of high-dimensional generative models following Ho et al. (2020) and Song et al. (2021). Their application to financial data is growing (Takahashi & Mizuno, 2025; Aghapour et al., 2025): denoising score matching on return panels can partially recover empirical stylized facts including excess kurtosis and cross-sectional correlation

structure. However, a *distributionally realistic* generator is not automatically a *portfolio-useful* one: the score model may generate return moments that systematically bias MVO weights, producing portfolios that are too concentrated, too volatile, or poorly calibrated to tail risk.

Two complementary tools address this gap. **Optimal transport (OT) calibration** aligns generated moments to historical moments post-hoc. The Gaussian Bures-Wasserstein map provides a closed-form linear correction matching means and covariances under a Gaussian approximation (Bures, 1969; Brenier, 1991; Dowson & Landau, 1982; Bhatia et al., 2019). **Stochastic control fine-tuning** directly reshapes the generative model to produce scenarios that improve a downstream portfolio objective. This second approach is grounded in KL-penalized stochastic optimal control, connecting the HJB equation to the Girsanov-cost control framework of Han et al. (2024); Fleming & Soner (2006); Øksendal (2003).

The tension is the trade-off between *realism* (staying close to the base generative model in KL sense) and *performance* (biasing scenarios toward high-utility portfolio outcomes). The KL penalty η parameterizes this frontier: smaller η permits larger drift perturbations, while larger η keeps the fine-tuned path measure closer to the base diffusion.

Contributions. (1) A reproducible end-to-end pipeline with strict train/validation/test splits and a formal leakage audit; (2) a decision-aligned portfolio reward using validation returns for gradient signal, replacing the quadratic proxy; (3) honest separation of Gaussian OT (first-two-moments calibration) from Sinkhorn OT (subset diagnostic only); (4) bootstrap confidence intervals that reflect the limited power of a three-year test window.




Figure 1: Method flow. Training data fit the scaler, score model, and OT target; validation data select η and supply the differentiable reward; test data are used only once after all choices are frozen.

2 Mathematical Framework

2.1 VP-SDE Score Model

Let $X_0 \sim p_{\text{data}}$ be a d -dimensional daily log-return vector. The variance-preserving forward process (Song et al., 2021) is:

$$dX_t = -\frac{1}{2}\beta(t)X_t dt + \sqrt{\beta(t)}dW_t, \quad \beta(t) = \beta_{\min} + t(\beta_{\max} - \beta_{\min}), \quad (1)$$

with $\beta_{\min} = 0.1, \beta_{\max} = 20$. The marginal distribution satisfies $X_t|X_0 \sim \mathcal{N}(\alpha_t X_0, \sigma_t^2 I_d)$, where $\alpha_t = e^{-\frac{1}{2}\int_0^t \beta(s)ds}$ and $\sigma_t^2 = 1 - \alpha_t^2$. At $t = 1$, $X_1 \approx \mathcal{N}(0, I_d)$ serves as the Gaussian prior. This affine Gaussian marginal follows by solving the linear SDE in (1) with an integrating factor and is the VP-SDE parameterization of Song et al. (2021).

A neural network $s_\theta(\mathbf{x}_t, t)$ is trained via *denoising score matching*:

$$\mathcal{L}_{\text{DSM}}(\theta) = \mathbb{E}_{t, X_0, \varepsilon} \left[\|s_\theta(\alpha_t X_0 + \sigma_t \varepsilon, t) + \varepsilon / \sigma_t\|^2 \right], \quad \varepsilon \sim \mathcal{N}(0, I_d). \quad (2)$$

The target is the conditional score $\nabla_{\mathbf{x}_t} \log p(\mathbf{x}_t|X_0) = -(\mathbf{x}_t - \alpha_t X_0)/\sigma_t^2 = -\varepsilon/\sigma_t$, which is the denoising score matching identity in Song & Ermon (2019); Song et al. (2021). Scenarios are generated via the Euler-Maruyama discretization of the reverse-time SDE (Anderson, 1982; Song et al., 2021). Euler-Maruyama is used because the learned score makes the reverse drift nonlinear and

unavailable in closed form; it is the standard first-order SDE integrator, with strong order 1/2, and is the baseline sampler in score-SDE implementations (Kloeden & Platen, 1992). In the implementation convention t_k decreases from 1 to 0 and $\Delta = 1/K$:

$$\mathbf{y}_{k+1} = \mathbf{y}_k + \left[-\frac{1}{2}\beta(t_k)\mathbf{y}_k + \beta(t_k)s_\theta(\mathbf{y}_k, t_k) \right] \Delta + \sqrt{\beta(t_k)\Delta} Z_k, \quad Z_k \sim \mathcal{N}(0, I_d). \quad (3)$$

2.2 HJB Formulation and KL-Penalized Fine-Tuning

We add a learnable drift perturbation $\mathbf{u}_\phi(\mathbf{y}, t)$ to the reverse SDE:

$$d\mathbf{y}_t = \left[-\frac{1}{2}\beta(t)\mathbf{y}_t + \beta(t)s_\theta(\mathbf{y}_t, t) + \mathbf{u}_\phi(\mathbf{y}_t, t) \right] dt + \sqrt{\beta(t)} dW_t. \quad (4)$$

Let p_ϕ be the path measure of (4) and p_θ the uncontrolled measure. By Girsanov’s theorem (Øksendal, 2003):

$$\text{KL}(p_\phi \| p_\theta) = \mathbb{E}_{\mathbf{y} \sim p_\phi} \left[\int_0^T \frac{\|\mathbf{u}_\phi(\mathbf{y}_t, t)\|^2}{2\beta(T-t)} dt \right]. \quad (5)$$

The fine-tuning objective is:

$$\max_{\phi} \mathbb{E}_{\mathbf{y}_0 \sim p_\phi} [R(\mathbf{y}_0)] - \eta \cdot \text{KL}(p_\phi \| p_\theta), \quad (6)$$

where $R : \mathbb{R}^d \rightarrow \mathbb{R}$ is a portfolio reward and $\eta > 0$ is the KL penalty. Writing $\mathcal{L}^{s_\theta}$ for the generator of the uncontrolled reverse sampler, the dynamic programming equation is

$$\partial_t V + \mathcal{L}^{s_\theta} V + \sup_{\mathbf{u}} \left\{ \mathbf{u}^\top \nabla_{\mathbf{y}} V - \frac{\eta}{2\beta(t)} \|\mathbf{u}\|^2 \right\} = 0, \quad V(\mathbf{y}, T) = R(\mathbf{y}). \quad (7)$$

The first-order condition gives

$$\mathbf{u}^*(\mathbf{y}, t) = \frac{\beta(t)}{\eta} \nabla_{\mathbf{y}} V(\mathbf{y}, t), \quad \partial_t V + \mathcal{L}^{s_\theta} V + \frac{\beta(t)}{2\eta} \|\nabla_{\mathbf{y}} V\|^2 = 0. \quad (8)$$

Equations (5)–(8) are the standard Girsanov/HJB derivation for drift control with quadratic relative-entropy cost (Fleming & Soner, 2006; Øksendal, 2003; Levine, 2018).

2.3 Gaussian OT Calibration

Given empirical moments $(\boldsymbol{\mu}_{\text{gen}}, \boldsymbol{\Sigma}_{\text{gen}})$ and target moments $(\boldsymbol{\mu}_{\text{train}}, \boldsymbol{\Sigma}_{\text{train}})$, the Bures-Wasserstein optimal map is (Bhatia et al., 2019):

$$T^*(\mathbf{x}) = \boldsymbol{\mu}_{\text{train}} + A(\mathbf{x} - \boldsymbol{\mu}_{\text{gen}}), \quad A = \boldsymbol{\Sigma}_{\text{gen}}^{-1/2} \left(\boldsymbol{\Sigma}_{\text{gen}}^{1/2} \boldsymbol{\Sigma}_{\text{train}} \boldsymbol{\Sigma}_{\text{gen}}^{1/2} \right)^{1/2} \boldsymbol{\Sigma}_{\text{gen}}^{-1/2}. \quad (9)$$

This map is used because it gives the exact W_2 -optimal affine coupling between Gaussian approximations to the generated and training distributions, so it directly corrects the mean and covariance errors that most affect MVO weights. Non-Gaussian structure (fat tails, volatility clustering) is *not* corrected by this linear map, which is why it is reported separately from Sinkhorn diagnostics (Peyré & Cuturi, 2019).

3 Empirical Design

3.1 Data and Splits

We use daily log-returns for ten S&P 500 sector ETFs (XLK, XLF, XLE, XLV, XLI, XLY, XLP, XLU, XLB, XLRE) via `yfinance` (Aroussi, 2019). The strict three-way split is:

Split	Dates	Trading Days	Use
Train	2014-01-01 to 2020-12-31	1,317	Score model, scaler, OT target
Validation	2021-01-01 to 2021-12-31	252	Reward eval, η selection
Test	2022-01-01 to 2024-12-31	756	Final evaluation (once)

A per-asset standardization scaler is fitted on training data only. Before final results are reported, an automated leakage check verifies that training, validation, and test information remain separated throughout scaling, model training, calibration, model selection, and final evaluation.

3.2 Score Model

A three-hidden-layer MLP score network ($d = 10$, hidden size 256, SiLU activations, sinusoidal time embedding $d_t = 64$) is trained with Adam (Kingma & Ba, 2015) at learning rate 10^{-3} , cosine annealing (Loshchilov & Hutter, 2017), 2000 epochs, batch size 256.

3.3 Portfolio Validation Reward

For a batch of generated scenarios $Y_0 \in \mathbb{R}^{B \times d}$, let $\hat{\mu} = \frac{1}{B} \sum_b Y_0^{(b)}$ and $\hat{\Sigma} = \text{Cov}(Y_0)$. The differentiable soft-Markowitz weight vector is:

$$\hat{\pi} = \text{softmax}\left((\hat{\Sigma} + \delta I)^{-1} \hat{\mu}\right), \quad \delta = 10^{-4}, \quad (10)$$

a differentiable approximation to the long-only MVO solution. The softmax is used because hard constrained MVO is not directly differentiable through the active set; it keeps weights positive, sums them to one, and permits gradient flow from the validation reward to generated scenario moments. Applied to fixed 2021 validation returns $r_{\text{val}} \in \mathbb{R}^{N_{\text{val}} \times d}$:

$$R(Y_0) = 252 \bar{r}_{\text{val}} - \frac{\lambda}{2} \cdot 252 \text{Var}(r_{\text{val}}) - \gamma \text{CVaR}_{95}(r_{\text{val}}) - \tau \text{HHI}(\hat{\pi}), \quad (11)$$

with $\lambda = 5$, $\gamma = 0.5$, $\tau = 0.1$, where $r_{\text{val}} = r_{\text{val}} \hat{\pi}$ and $\text{HHI}(\hat{\pi}) = \sum_i \hat{\pi}_i^2$ is the Herfindahl concentration index. Gradient flows from r_{val} through $\hat{\pi}$ to $(\hat{\mu}, \hat{\Sigma})$ to the generated samples Y_0 . CVaR is included because downside tail losses are central in portfolio evaluation and are not captured by variance alone (Rockafellar & Uryasev, 2000; McNeil et al., 2015); HHI discourages degenerate single-sector allocations.

3.4 η Selection Protocol

The control network u_ϕ (tanh-bounded MLP, hidden 128) is trained for 200 epochs for each candidate $\eta \in \{0.05, 0.10, 0.50, 1.0, 5.0\}$. The selection criterion is the 2021 *validation* Sharpe ratio of implied MVO weights. The selected η^* is fixed before any 2022–2024 test-period evaluation is run.

4 Results

4.1 Leakage Audit

The leakage audit checks whether any test-period information could have entered model training, calibration, or hyperparameter selection. For the run reported here, all eight checks passed with no failures or warnings: train, validation, and test dates do not overlap; scaling, score-model training, and Gaussian OT calibration use only 2014–2020 training data; η selection uses only the 2021 validation Sharpe; final test metrics are computed only after η is frozen; and generated scenario NaN/Inf rates are below 1%. This result does not certify that the strategy will generalize; it certifies that the reported 2022–2024 test backtest was not used to tune the model choices.

4.2 Base Scenario Quality

Figure 2 compares generated and historical returns across stylized facts. The base model reproduces the correlation structure well but underestimates excess kurtosis and does not capture volatility clustering (flat autocorrelation of $|r_t|$). Table 1 summarizes OT calibration diagnostics.

Table 1: OT calibration diagnostics: Bures-Wasserstein distance to training data. Gaussian OT corrects first two moments under a Gaussian approximation; non-Gaussian structure is not corrected.

Scenario set	W_2 to train	HHI (implied MVO)
Base generated	0.0303	0.629
After Gaussian OT	0.0003	0.260
Sinkhorn (subset)	0.0205	—

Sinkhorn applied to 500-sample subset only; not a full portfolio strategy.




Figure 2: Stylized fact comparison between generated scenarios (orange) and 2014–2020 historical returns (blue). Top left: excess kurtosis by asset. Top center: autocorrelation of $|r_t|$ (volatility clustering). Top right: historical correlation matrix. Bottom left: generated correlation matrix. Bottom center: marginal return distribution (XLK). Bottom right: QQ-plot (XLK). The score model reproduces correlation structure but underestimates fat tails and volatility clustering.

4.3 OT Calibration

Figure 3 shows the effect of Gaussian OT on per-asset volatility and correlation structure.




Figure 3: Gaussian OT diagnostics: W2 distance, marginal distributions, per-asset standard deviations, and correlation matrices before/after calibration. The linear map A corrects both mean and covariance under the Gaussian approximation.

4.4 η Selection (Validation Only)

Table 2 reports validation metrics for each η candidate. The selected $\eta^* = 0.5$ achieves the highest validation Sharpe (2.30). Small η (over-steering) and large η (insufficient control) both degrade validation performance.

Table 2: η sweep: validation period (2021) metrics. No test metrics are available at selection time. $\eta^* = 0.5$ selected by validation Sharpe.

η	Val. Sharpe	Val. Ann. Return	Val. CVaR(95%)	Val. HHI
0.05	0.97	—	—	—
0.10	1.95	21.85%	0.0174	0.479
0.50	2.30	34.93%	0.0203	0.316
1.00	1.82	33.66%	0.0257	0.738
5.00	2.14	32.39%	0.0205	0.334




Figure 4: η selection frontier based on 2021 validation data only. Left: validation Sharpe vs. η (log scale). Right: validation CVaR95 vs. η . The star marks the selected $\eta^* = 0.5$. No test-period information enters this selection.

4.5 Fixed-Scenario Backtest (Test 2022–2024)

Table 3 and Figure 5 report final test metrics. The fine-tuned strategy achieves the highest Sharpe (0.70), highest return (11.31%), and lowest maximum drawdown (14.88%). Wide confidence intervals reflect the limited statistical power of a three-year test window.

Table 3: Fixed-scenario backtest results, test period 2022–2024. Evaluated once after η frozen. Bootstrap CI: 2000 draws, i.i.d. (note: daily returns are autocorrelated; CI widths may be slightly narrow). No model selection occurs here.

Strategy	Ann. Return	Ann. Vol.	Sharpe	95% CI	CVaR (95%)	Max DD	HHI
Equal-Weight	5.59%	15.41%	0.363	$[-0.76, +1.46]$	0.0226	−18.4%	0.100
Markowitz	9.03%	17.54%	0.515	$[-0.60, +1.57]$	0.0264	−17.1%	0.541
Wasserstein-Robust	5.75%	15.12%	0.380	$[-0.74, +1.48]$	0.0222	−17.9%	0.103
Diffusion-Base	8.69%	22.53%	0.386	$[-0.72, +1.49]$	0.0319	−30.7%	0.629
Diffusion+GaussOT	6.45%	16.56%	0.390	$[-0.69, +1.49]$	0.0242	−23.0%	0.260
Diffusion+FT($\eta^*=0.5$)	11.31%	16.16%	0.700	$[-0.38, +1.76]$	0.0241	−14.9%	0.316

Note: Confidence intervals span ≈ 2.1 Sharpe units at 95% level. No pairwise difference is individually statistically significant. Results are exploratory evidence of a realism-performance frontier.




Figure 5: Fixed-scenario backtest (2022–2024). Left: cumulative wealth. Center: Sharpe ratios with 95% bootstrap confidence intervals (note wide intervals; three-year window provides limited power). Right: drawdown profiles. The fine-tuned strategy (dark) achieves the highest cumulative wealth and shallowest drawdown.




Figure 6: Drawdown profiles in the fixed-scenario backtest. Diffusion-Base has the deepest drawdown (−30.7%), while Gaussian OT and fine-tuning both reduce drawdown relative to the uncalibrated base generator.

4.6 Rolling-Diffusion Backtest (Test 2022–2024)

Table 4: Rolling-diffusion backtest results, test period 2022–2024. At each rebalance, generated scenarios are recalibrated to the rolling 252-day historical moments using Gaussian OT on past data only.

Strategy	Sharpe	Ann. Return	Max DD	HHI	Turnover
Equal-Weight	0.363	5.59%	−18.4%	0.100	0.000
Markowitz	0.515	9.03%	−17.1%	0.541	0.617
Wasserstein-Robust	0.380	5.75%	−17.9%	0.103	0.040
Diffusion-Base	0.510	8.89%	−17.1%	0.496	0.585
Diffusion+GaussOT	0.510	8.87%	−17.1%	0.498	0.586
Diffusion+FT(η^*)	0.510	8.88%	−17.1%	0.496	0.585

Rolling recalibration causes all diffusion strategies to match rolling Markowitz performance (Table 4). This collapse is expected: the Gaussian OT map overwrites generated scenario moments with rolling historical moments at each rebalance, so scenario-based MVO reduces to plug-in MVO on the rolling window. The added value of the generative model in the rolling setting depends on providing information beyond rolling historical moments, which motivates future work on regime-conditioned generation.

5 Discussion

Realism-performance trade-off. The main finding is that KL-penalized fine-tuning creates a meaningful realism-performance frontier. At $\eta^* = 0.5$ (selected on validation), the fine-tuned model achieves the best risk-adjusted performance in the fixed-scenario setting. At smaller η , the model over-steers the distribution and generates out-of-distribution scenarios that degrade validation performance. At larger η , the control is insufficient to materially differentiate the portfolio allocation.

Why rolling mode collapses. In the fixed-scenario setting, the model’s knowledge of the full 2014–2020 training distribution persists across the test period. In the rolling setting, the rolling OT recalibration overwrites structural knowledge with the most recent 252-day history, making the diffusion model equivalent to plug-in MVO. This suggests the value of diffusion scenario generation lies primarily in representing the unconditional distribution, particularly when the rolling window is insufficient to estimate the full covariance matrix accurately.

Statistical caution. A three-year test window provides ≈ 756 observations. At this sample size, 95% bootstrap Sharpe CI width is ≈ 2.1 Sharpe units (Table 3). Most pairwise differences are within one standard error. Results must be interpreted as exploratory, consistent with prior work in this horizon (Bailey & Lopez de Prado, 2014).

6 Limitations and Future Work

Score model expressiveness. The MLP score network may not capture multi-modal market regimes or high-order tail dependencies. Transformer-based score functions (Dhariwal & Nichol, 2021) or flow matching (Lipman et al., 2023) may improve distributional realism.

Gaussian OT scope. The Bures-Wasserstein map corrects only first and second moments. Full Sinkhorn OT (Peyré & Cuturi, 2019) or neural OT (Bunne et al., 2022) would address higher-order structure at greater computational cost.

Differentiable reward. The soft-Markowitz approximation (10) replaces constrained MVO with a softmax projection. Differentiable convex programming layers (Agrawal et al., 2019) would give exact long-only weights while preserving gradient flow.

Bootstrap CI. Simple i.i.d. bootstrap underestimates CI width for autocorrelated returns; block bootstrap (Politis & Romano, 1994) would give wider, more appropriate intervals.

Scalability. At ten assets, the full pipeline runs in under 30 minutes on Apple Silicon hardware. Scaling to hundreds of assets would require mini-batch Sinkhorn (Favras et al., 2021) and factored covariance approximations.

7 Conclusion

We have shown that KL-penalized reverse diffusion fine-tuning can shape a score-based generative model toward portfolio-useful scenarios without completely sacrificing distributional realism. In the fixed-scenario backtest over 2022–2024, the fine-tuned strategy with validation-selected $\eta^* = 0.5$ achieves Sharpe 0.70, annualized return 11.31%, and maximum drawdown 14.88%, compared to 0.515 Sharpe and 17.09% drawdown for rolling Markowitz.

Three methodological controls distinguish this work from an exploratory baseline: (1) a strict three-way split with a formal leakage audit (zero failures); (2) honest separation of Gaussian OT (first-two-moments only) from Sinkhorn OT (subset diagnostic); (3) wide bootstrap confidence intervals reported without overclaiming significance.

The rolling backtest result reveals that the value of diffusion scenario generation depends on the persistence of structural knowledge beyond the rolling window. Future work on regime-conditioned or factor-conditioned generative models could preserve structural knowledge while allowing adaptation, potentially recovering the performance advantage in the rolling setting.

References

- Aghapour, A., Bayraktar, E., & Yuan, F. (2025). Solving dynamic portfolio selection problems via score-based diffusion models. *arXiv preprint arXiv:2507.09916*.
- Agrawal, A., Amos, B., Barratt, S., Boyd, S., Diamond, S., & Kolter, J. Z. (2019). Differentiable convex optimization layers. *Advances in Neural Information Processing Systems*, 32, 9562–9574.
- Anderson, B. D. O. (1982). Reverse-time diffusion equation models. *Stochastic Processes and Their Applications*, 12(3), 313–326.

- Aroussi, R. (2019). *yfinance: Yahoo! Finance market data downloader* [Software]. <https://github.com/ranaroussi/yfinance>.
- Bailey, D. H., & Lopez de Prado, M. (2014). The deflated Sharpe ratio: Correcting for selection bias, backtest overfitting, and non-normality. *Journal of Portfolio Management*, 40(5), 94–107.
- Bhatia, R., Jain, T., & Lim, Y. (2019). On the Bures-Wasserstein distance between positive definite matrices. *Expositiones Mathematicae*, 37(2), 165–191.
- Brenier, Y. (1991). Polar factorization and monotone rearrangement of vector-valued functions. *Communications on Pure and Applied Mathematics*, 44(4), 375–417.
- Bunne, C., Krause, A., & Cuturi, M. (2022). Supervised training of conditional Monge maps. *Advances in Neural Information Processing Systems*, 35, 6859–6872.
- Bures, D. (1969). An extension of Kakutani’s theorem on infinite product measures to the tensor product of semifinite W^* -algebras. *Transactions of the American Mathematical Society*, 135, 199–212.
- Cont, R. (2001). Empirical properties of asset returns: Stylized facts and statistical issues. *Quantitative Finance*, 1(2), 223–236.
- Dhariwal, P., & Nichol, A. (2021). Diffusion models beat GANs on image synthesis. *Advances in Neural Information Processing Systems*, 34, 8780–8794.
- Dowson, D. C., & Landau, B. V. (1982). The Fréchet distance between multivariate normal distributions. *Journal of Multivariate Analysis*, 12(3), 450–455.
- Fatras, K., Sejourne, T., Flamary, R., & Courty, N. (2021). Unbalanced minibatch optimal transport: Applications to domain adaptation. *Proceedings of the 38th ICML*, 3186–3197.
- Fleming, W. H., & Soner, H. M. (2006). *Controlled Markov Processes and Viscosity Solutions* (2nd ed.). Springer.
- Han, Y., Razaviyayn, M., & Xu, R. (2024). Stochastic control for fine-tuning diffusion models: Optimality, regularity, and convergence. *arXiv preprint arXiv:2412.18164*.
- Ho, J., Jain, A., & Abbeel, P. (2020). Denoising diffusion probabilistic models. *Advances in Neural Information Processing Systems*, 33, 6840–6851.
- Kingma, D. P., & Ba, J. (2015). Adam: A method for stochastic optimization. *Proceedings of ICLR 2015*.
- Kloeden, P. E., & Platen, E. (1992). *Numerical Solution of Stochastic Differential Equations*. Springer.
- Levine, S. (2018). Reinforcement learning and control as probabilistic inference: Tutorial and review. *arXiv preprint arXiv:1805.00909*.
- Lipman, Y., Chen, R. T. Q., Ben-Hamu, H., Nickel, M., & Le, M. (2023). Flow matching for generative modeling. *Proceedings of ICLR 2023*.
- Loshchilov, I., & Hutter, F. (2017). SGDR: Stochastic gradient descent with warm restarts. *Proceedings of ICLR 2017*.
- McNeil, A. J., Frey, R., & Embrechts, P. (2015). *Quantitative Risk Management: Concepts, Techniques and Tools* (rev. ed.). Princeton University Press.
- Markowitz, H. (1952). Portfolio selection. *Journal of Finance*, 7(1), 77–91.

- Øksendal, B. (2003). *Stochastic Differential Equations: An Introduction with Applications* (6th ed.). Springer.
- Peyré, G., & Cuturi, M. (2019). Computational optimal transport. *Foundations and Trends in Machine Learning*, 11(5–6), 355–607.
- Politis, D. N., & Romano, J. P. (1994). The stationary bootstrap. *Journal of the American Statistical Association*, 89(428), 1303–1313.
- Rockafellar, R. T., & Uryasev, S. (2000). Optimization of conditional value-at-risk. *Journal of Risk*, 2(3), 21–41.
- Song, Y., & Ermon, S. (2019). Generative modeling by estimating gradients of the data distribution. *Advances in Neural Information Processing Systems*, 32, 11895–11907.
- Song, Y., Sohl-Dickstein, J., Kingma, D. P., Kumar, A., Ermon, S., & Poole, B. (2021). Score-based generative modeling through stochastic differential equations. *Proceedings of ICLR 2021*.
- Takahashi, T., & Mizuno, T. (2025). Generation of synthetic financial time series by diffusion models. *Quantitative Finance*, 25(10), 1507–1516.

A Reproducibility

The pipeline is reproducible as a fixed sequence of stages: data download and splitting, leakage audit, score-model training, scenario generation, OT calibration, validation-only η selection, one-shot test evaluation, and figure/table generation. The production run uses 2,000 score-model training epochs and 10,000 generated scenarios; a reduced smoke-test setting is available for checking the full workflow quickly. Random seeds are fixed at 42 for data and score-model training and 0 for fine-tuning.

B Leakage Audit Details

The leakage audit checks the following conditions and passed all of them in the reported run:

1. Date ranges of train, validation, and test windows do not overlap.
2. The scaler is fitted on training data only.
3. The score model is trained on 2014–2020 data with no test-period reference.
4. Gaussian OT target moments are computed from training data only.
5. Sinkhorn OT labeled as subset-only; test period not referenced.
6. Validation Sharpe is the only criterion used for η selection.
7. Final metrics are generated only after the validation-selected η has been frozen, so no test-period result is used for model selection.
8. Scenario NaN and Inf fractions below 0.01.

The observed result for the report run is pass: zero failures and zero warnings.

C Fine-Tuning Training Curves




Figure 7: Fine-tuning training curves. Left: portfolio reward (Equation 11) across epochs; reward increases monotonically, confirming the control successfully shapes the distribution. Right: KL cost (Equation 5) per epoch; the control effort stabilizes after approximately 50 epochs.